

# Google Agent Development Kit: A Comprehensive Overview

## Executive Summary:

The Google Agent Development Kit (ADK), introduced at Google Cloud NEXT 2025, marks a significant advancement in the field of artificial intelligence by providing a robust open-source framework for the streamlined creation of sophisticated AI agents, with a particular emphasis on multi-agent systems.<sup>1</sup> This toolkit empowers developers with the necessary flexibility and precise control to build agentic applications that are ready for production deployment.<sup>1</sup> The ADK's core strengths include its developer-centric "code-first" approach, its deep integration with the Google Cloud ecosystem, particularly with Gemini models and Vertex AI, and its extensive compatibility with a wide array of language models through the inclusion of LiteLLM.<sup>1</sup> Furthermore, the ADK offers a rich set of tools that allow for the extension of agent capabilities, along with inherent support for features such as real-time streaming, thorough performance evaluation, and the development of responsible AI applications.<sup>1</sup> The primary aim of the ADK is to tackle the inherent complexities associated with developing multi-agent systems by offering developers precise control over agent behavior and the orchestration of their interactions, a seamlessly integrated development experience, and a comprehensive framework for evaluating agent performance.<sup>1</sup> The open-source nature of the ADK is intended to foster innovation and encourage contributions from the broader developer community within the rapidly evolving landscape of AI agents.<sup>1</sup>

The central design principle of the ADK is the facilitation of multi-agent systems. This indicates a forward-thinking strategy that acknowledges the increasing complexity of AI applications and the potential of collaborative intelligence. Multiple sources highlight the ADK's focus on enabling the creation of systems where multiple specialized agents can work together to achieve complex goals.<sup>1</sup> This emphasis suggests that Google anticipates the future of advanced AI applications will rely on the coordinated efforts of various intelligent entities, and the ADK is specifically architected to support this development paradigm.

Another significant aspect of the ADK is its "code-first" philosophy. This approach, repeatedly mentioned in the documentation and related articles, signifies a deliberate focus on developers who desire a high degree of control and customization over their AI agents.<sup>3</sup> By requiring developers to define agent logic and behavior directly through Python code, the ADK caters to those who need fine-grained control and the ability to implement intricate functionalities, contrasting with more abstract, low-code or no-code solutions.

Furthermore, the ADK's tight integration with Google Cloud services is a key characteristic that positions it as a valuable tool for developers operating within this ecosystem.<sup>1</sup> Several sources point out the ADK's optimization for Google Cloud, its seamless compatibility with Gemini models, and its deployment capabilities on Vertex

AI. This close relationship suggests that developers already utilizing or planning to utilize Google Cloud resources will find the ADK particularly advantageous due to the optimized performance and streamlined deployment options it provides.

## **1. Introduction to the Google Agent Development Kit (ADK):**

### **1.1 Core Purpose:**

The Agent Development Kit (ADK) is specifically designed to empower developers in the multifaceted process of building, managing, evaluating, and deploying sophisticated AI-powered agents.<sup>7</sup> This encompasses a broad spectrum of agent types, including conversational agents capable of engaging in natural language dialogue with users, as well as non-conversational agents designed to autonomously execute specific tasks and workflows.<sup>7</sup> A primary objective of the ADK is to streamline the entire end-to-end development lifecycle for both individual agents and intricate multi-agent systems.<sup>1</sup> This simplification covers all stages, from the initial conceptualization and implementation of an agent to its rigorous testing, efficient deployment into production environments, and ongoing management and maintenance.<sup>1</sup> The ADK provides a robust and highly adaptable environment that enables the creation of AI agents capable of addressing complex tasks and orchestrating intricate workflows.<sup>7</sup> This inherent flexibility allows developers to tailor agents to a wide variety of applications across diverse domains, ensuring that the framework can accommodate a broad range of use cases. The design of the ADK, with its emphasis on providing developers with a flexible yet powerful environment for building complex AI applications, suggests that its intended users are experienced software engineers and AI specialists who require a high degree of control over their agent development process. The terminology used, such as "empower developers," "robust and flexible environment," and "complex tasks and workflows," indicates that the ADK is not geared towards novice users or those seeking simplistic, out-of-the-box solutions. Instead, it is intended for professionals who need a comprehensive toolkit to construct advanced AI applications tailored to specific and often demanding requirements.

### **1.2 Key Benefits:**

The ADK offers developers a significant advantage by providing precise control over how their agents behave and how interactions are orchestrated within multi-agent systems.<sup>1</sup> This level of granular control is essential for building AI applications that are not only intelligent but also reliable and predictable in their operation.<sup>1</sup> By its very design, the ADK facilitates the creation of applications that are both modular and highly scalable. This is achieved through the ability to compose multiple specialized agents into a hierarchical structure.<sup>1</sup> This architectural approach allows for the development of complex systems where different agents are responsible for specific functionalities, thereby promoting better code organization, easier maintenance, and the potential for reusing agent components across different applications.<sup>1</sup> A notable feature of the ADK is its built-in support for bidirectional audio and video streaming capabilities.<sup>1</sup> This enables developers to create more natural and human-like conversational interfaces with relative ease, often requiring just a few lines of code.<sup>1</sup> This capability moves beyond traditional text-based interactions, allowing for the development of richer, multimodal

dialogues that can enhance user experience and expand the possibilities for agent interaction.<sup>1</sup> The framework also boasts seamless integration with a rich ecosystem of language models. This includes Google's advanced Gemini models, which offer cutting-edge performance in various AI tasks.<sup>1</sup> Additionally, the ADK supports any model accessible through the Vertex AI Model Garden, providing a wide selection of pre-trained models from Google and other leading AI providers.<sup>1</sup> Furthermore, through its integration with LiteLLM, the ADK extends its compatibility to include a vast range of third-party models from providers such as Anthropic, Meta, Mistral AI, and AI21 Labs.<sup>1</sup> This extensive model support gives developers the flexibility to choose the most appropriate model for their specific needs and use cases.<sup>1</sup> Developers gain access to a diverse array of tools within the ADK, further enhancing the capabilities of their agents.<sup>2</sup> These include function tools for integrating custom Python code, built-in tools for common functionalities like web search and code execution, third-party tool integrations such as LangChain and CrewAI, Google Cloud tools for leveraging various cloud services, MCP tools for connecting to diverse data sources, and OpenAPI tools for interacting with external APIs.<sup>2</sup> The ADK significantly streamlines the process of deploying agents into production environments through its integration with Vertex AI Agent Engine, a fully managed, scalable, and enterprise-grade runtime.<sup>1</sup> It also supports deployment on Google Cloud Run and other container runtimes, offering flexibility in choosing the deployment environment.<sup>1</sup> This simplified deployment process allows for a smoother transition from development to production.<sup>1</sup> The ADK provides an integrated developer experience that aims to simplify the entire process of building, testing, and debugging AI agents.<sup>1</sup> This includes various tools and interfaces designed to make the development workflow more efficient and intuitive.<sup>1</sup> A robust evaluation framework is also built into the ADK, enabling developers to systematically assess the performance and effectiveness of their agents.<sup>1</sup> This evaluation process includes analyzing both the quality of the final output generated by the agent and the individual steps taken by the agent to reach that output, providing a comprehensive understanding of its performance.<sup>1</sup> The ADK offers native support for bidirectional streaming of both text and audio, facilitating the creation of real-time, interactive experiences with AI agents.<sup>5</sup> This feature is particularly valuable for building conversational agents that can provide immediate feedback and engage in more natural interactions.<sup>1</sup> The framework also includes features for managing short-term conversational context (session), maintaining the agent's state across interactions, handling long-term memory to retain information over time, and managing artifacts such as file uploads and downloads.<sup>5</sup> Furthermore, the ADK is designed to be highly extensible, allowing developers to deeply customize agent behavior using callbacks and easily integrate third-party tools and services to meet specific application requirements.<sup>5</sup> The sheer breadth of these benefits, ranging from the ability to precisely control agent behavior to the readiness for production deployment, indicates that the ADK is a comprehensive solution designed to address the entire lifecycle of AI agent development. It is not merely focused on a single aspect, such as model integration or deployment, but rather aims to provide a holistic toolkit for building and managing AI agents across various complexities and deployment scenarios.

### 1.3 Target Users:

The ADK is primarily intended for developers who are seeking to build agentic applications

that are production-ready and capable of handling real-world use cases.<sup>1</sup> It is particularly well-suited for developers who require a high degree of flexibility and precise control over the behavior and orchestration of their AI agents.<sup>5</sup> The framework is designed to cater to both individual developers and teams working on complex AI agents and sophisticated multi-agent systems.<sup>5</sup> Moreover, developers who plan to leverage the powerful AI models and robust infrastructure offered by the Google Cloud ecosystem will find the ADK especially advantageous due to its deep and seamless integration with these services.<sup>1</sup>

## **2. Getting Started with ADK:**

### **2.1 Installation:**

The primary and recommended method for installing the Agent Development Kit (ADK) is through the use of pip, the Python package installer.<sup>7</sup> Developers can easily install the core ADK package by executing the command `pip install google-adk` in their terminal.<sup>7</sup> To ensure optimal compatibility and access to all the features of the ADK, developers need to have Python version 3.10 or a more recent version installed on their development machine.<sup>10</sup> Utilizing a supported Python version guarantees that the framework's functionalities and dependencies will operate as expected.<sup>10</sup> It is strongly advised that developers create and activate a virtual environment before proceeding with the installation of the ADK.<sup>10</sup> This practice is crucial for isolating the dependencies of the specific ADK project from other Python projects on the system.<sup>10</sup> By creating a dedicated virtual environment, developers can prevent potential conflicts between different project dependencies and maintain a clean and organized development environment.<sup>10</sup> For developers seeking more detailed and step-by-step guidance on the installation process, the official ADK documentation provides a comprehensive Installation guide that covers various aspects of setting up the development environment.<sup>7</sup>

### **2.2 Quickstart:**

The ADK is intentionally designed to enable developers to quickly begin building and experimenting with AI agents. It allows for the creation of a functional agent integrated with tools in a remarkably short timeframe, often within minutes.<sup>7</sup> This rapid setup capability is essential for developers who want to quickly explore the framework's potential and start prototyping their ideas.<sup>7</sup> The Quickstart guide offers a detailed walkthrough of the initial steps required to set up and run a basic agent using the ADK.<sup>7</sup> This guide covers essential aspects such as setting up the development environment, defining the agent's core logic, and initiating basic interactions with the agent.<sup>7</sup> The quickstart process typically involves establishing a specific folder structure for the agent project.<sup>11</sup> This structure usually includes a parent directory containing a subdirectory dedicated to the agent itself (for example, `multi_tool_agent`), along with several key files.<sup>11</sup> These files include `__init__.py`, which is necessary to make the agent directory a Python package, `agent.py`, which will house the main logic and definition of the agent, and `.env`, a file used to store environment-specific variables.<sup>11</sup> Within the `agent.py` file, developers will define the core intelligence and behavior of their AI agent.<sup>11</sup> The quickstart guide often provides illustrative example code that demonstrates how to create an agent with fundamental capabilities, such as the ability to retrieve weather information or the current time for a specified location using predefined

tools.<sup>11</sup> The `.env` file serves as a secure place to store sensitive information, such as API keys required to authenticate with the underlying Large Language Model (LLM) that powers the agent's ability to understand and respond to requests.<sup>11</sup> The quickstart guide offers clear instructions on how to configure this file with the necessary credentials, whether the developer chooses to use Gemini through Google AI Studio or Gemini through Vertex AI on Google Cloud.<sup>11</sup> The ADK provides developers with multiple convenient ways to run and test their newly created agent locally.<sup>11</sup> These methods include a developer Web UI, which can be launched using the command `adk web`, providing an interactive browser-based interface for interacting with the agent.<sup>11</sup> Alternatively, developers can interact with their agent directly from the terminal using the `adk run` command.<sup>11</sup> For those who prefer to test programmatic interactions, the `adk api_server` command starts a local FastAPI server, allowing for testing via cURL requests.<sup>11</sup> The quickstart guide typically includes a set of example prompts that developers can use to test the basic functionalities of their agent, such as asking for the weather in a particular city or the current time.<sup>10</sup>

### 2.3 Quickstart (streaming):

For developers who are interested in building AI agents capable of providing real-time, streaming responses, the ADK offers a dedicated quickstart guide specifically for creating streaming agents.<sup>7</sup> This capability allows for the development of more interactive and engaging user experiences, where the agent's response is generated and displayed incrementally as it is being processed.<sup>7</sup> More comprehensive information and detailed instructions on how to create streaming agents using the ADK can be found in the Quickstart (streaming) guide.<sup>7</sup> It is important to note that snippet 12 indicates that the URL for the Quickstart (streaming) guide is currently inaccessible. This suggests a potential issue with the documentation that should be taken into consideration.

### 2.4 Tutorial: Building a Multi-Agent System:

To effectively demonstrate the ADK's capabilities in handling more intricate and complex scenarios, the official documentation includes a comprehensive tutorial that guides developers through the process of building their first multi-agent system using the ADK.<sup>7</sup> This tutorial serves as a practical illustration of how to structure interactions between multiple independent AI "brains" that work collaboratively to achieve a common goal.<sup>7</sup> The tutorial specifically focuses on the creation of a "Weather Bot" team, which effectively showcases the concept of delegation, where different specialized agents within the system are responsible for handling specific parts of a user's request.<sup>8</sup> A crucial aspect of this multi-agent tutorial involves the definition of specialized tools that individual sub-agents can utilize to perform their designated tasks.<sup>8</sup> Examples provided include tools for generating greetings (`say_hello`), providing farewell messages (`say_goodbye`), and retrieving weather-related information (`get_weather`).<sup>8</sup> The tutorial then explains in detail how to define individual sub-agents, such as a `greeting_agent` and a `farewell_agent`, each with very specific instructions that outline their sole purpose and the exact tools they are authorized to use.<sup>8</sup> These highly focused instructions are essential for ensuring effective delegation of tasks within the multi-agent system.<sup>8</sup> Following the definition of the sub-agents, the tutorial guides developers in creating a root agent, which in this case is an updated version of the weather agent, specifically configured to utilize the previously defined sub-agents.<sup>8</sup> This is achieved through the use of

the `sub_agents` parameter, which allows the root agent to intelligently delegate specific types of user requests to the most appropriate sub-agent based on the descriptions provided for each sub-agent.<sup>8</sup> The tutorial then demonstrates how to interact with the complete multi-agent system by sending various types of queries.<sup>8</sup> It illustrates how the root agent can recognize a greeting based on the description of the `greeting_agent` and automatically delegate the request to it. Similarly, it shows how weather-related requests are handled directly by the root agent using the `get_weather` tool, and how farewell messages are delegated to the `farewell_agent`.<sup>8</sup> The tutorial also highlights the significant role of LiteLLM as a key component that enables multi-model support within the ADK.<sup>8</sup> This allows different agents within the same multi-agent system to potentially be powered by different underlying language models, offering greater flexibility and the ability to choose the best model for each specific task.<sup>8</sup> The detailed step-by-step instructions and comprehensive explanations for building this multi-agent system can be found in the(<https://google.github.io/adk-docs/get-started/tutorial/>) section of the official documentation.<sup>7</sup> The early introduction of a multi-agent tutorial in the "Getting Started" section emphasizes the ADK's core focus on building collaborative AI systems. By providing this tutorial early on, the ADK developers signal that multi-agent functionality is a fundamental capability of the framework and encourage users to consider designing their applications using a team of specialized agents working in concert.

### 2.5 Local Testing:

The ADK provides various resources and tools to facilitate thorough local testing of AI agents before they are deployed to production environments.<sup>7</sup> This capability is crucial for allowing developers to iterate quickly on their agent designs and ensure that their agents function as intended in a controlled and isolated environment.<sup>7</sup> As mentioned in the quickstart guide, the `adk web` command is used to launch a developer Web UI, which provides an interactive platform for testing agent behavior and inspecting the flow of interactions.<sup>11</sup> Similarly, the `adk run` command allows developers to test their agents directly from the command-line terminal.<sup>11</sup> For more in-depth information and advanced techniques for conducting local testing, developers can refer to the dedicated documentation on local testing provided by the ADK.<sup>11</sup>

### 2.6 Sample Agents:

To further aid developers in understanding and effectively utilizing the ADK, the documentation includes a valuable collection of sample agents designed for various real-world applications.<sup>7</sup> These practical examples cover diverse domains such as retail, travel, and customer service, offering concrete demonstrations of how to build agents tailored to specific use cases.<sup>7</sup> Developers can explore the code and configurations of these sample agents in the `adk-samples` repository hosted on GitHub.<sup>7</sup> This repository serves as an invaluable resource for learning best practices in agent development and discovering reusable components that can accelerate the development process.<sup>7</sup>

### 2.7 About ADK:

The "About ADK" section of the documentation provides a high-level overview of the key components and fundamental concepts that underpin the process of building and deploying AI agents using the framework.<sup>7</sup> This section is intended to help developers gain a

foundational understanding of the ADK's overall architecture and the core principles that guide its design.<sup>7</sup> It reinforces the fact that the ADK is a Python framework specifically engineered to streamline the development of applications powered by Large Language Models (LLMs).<sup>8</sup> The "About ADK" section emphasizes the framework's role in providing robust and well-defined building blocks that empower developers to create intelligent agents.<sup>8</sup> Furthermore, this section highlights the ADK's capabilities in enabling the creation of agents that can reason about information, formulate plans to achieve goals, effectively utilize external tools, interact dynamically with users through natural language, and collaborate seamlessly within a team of other agents.<sup>8</sup>

### **3. Deep Dive into Agent Architectures:**

#### **3.1 LLM Agents:**

At the core of the ADK's agent capabilities are LLM agents, which harness the power of Large Language Models to execute complex tasks.<sup>2</sup> These agents possess the ability to understand natural language input, reason about the information they process, make informed decisions based on the context of the interaction, and interact with various external tools to achieve their designated objectives.<sup>2</sup> The `LlmAgent` class, often conveniently aliased as simply `Agent`, serves as the fundamental "thinking" component within any ADK application.<sup>2</sup> Developers instantiate this class to create individual agents that are powered by the underlying LLM and then define their specific behaviors and inherent capabilities.<sup>2</sup> Constructing an effective LLM agent involves carefully defining three key components: Identity, Instruction, and Tools.<sup>2</sup> The Identity of an agent consists of a unique name and a clear, concise description.<sup>2</sup> This description is particularly important in multi-agent systems as it helps other agents in the system understand the specific expertise and capabilities of this agent, allowing them to make informed decisions about when to delegate particular tasks.<sup>2</sup> The Instruction is a critical string of text that fundamentally shapes the agent's behavior.<sup>2</sup> This instruction defines the agent's primary tasks, its intended personality (if applicable to the application), any constraints or limitations it must adhere to, and crucially, how it should effectively utilize the available tools to accomplish its goals.<sup>2</sup> The instruction acts as the primary driver for the LLM's reasoning processes and its decision-making during interactions.<sup>2</sup> Tools represent the external capabilities that extend the agent's functionality beyond the inherent knowledge contained within the LLM itself.<sup>2</sup> These tools allow the agent to interact with the real world, access up-to-date information from external sources, or perform specific actions that are necessary to fulfill its purpose.<sup>2</sup> LLM agents offer a significant degree of flexibility in how their behavior is defined, primarily through the instruction parameter.<sup>2</sup> This allows developers to create agents with diverse personalities, ranging from highly formal and task-oriented to more conversational and engaging, and to imbue them with specialized skills tailored to specific application domains.<sup>2</sup> In the context of multi-agent scenarios, LLM agents can be designed to intelligently delegate tasks to other agents within the system.<sup>1</sup> This delegation is typically based on the current context of the conversation and the advertised capabilities of the other agents, as described in their respective identities.<sup>1</sup> This enables the creation of complex workflows where different agents are responsible for handling specific parts of a larger, overarching task.<sup>1</sup> The ADK also supports LLM-driven dynamic routing.<sup>4</sup> In this paradigm, the

LLM agent itself makes the decision about which tool to utilize or which sub-agent to delegate a task to, based on the user's input and the current state of the interaction.<sup>4</sup> This approach allows for more adaptive and contextually aware agent behavior, as the agent can dynamically adjust its strategy based on the evolving needs of the interaction.<sup>4</sup> LLM agents within the ADK can be extensively customized by carefully crafting their instructions, providing them with access to a specific and relevant set of tools, and integrating them into various types of workflows and more complex multi-agent architectures.<sup>3</sup> The LlmAgent class provides a powerful and intuitive abstraction for leveraging the advanced capabilities of large language models within the ADK. The clear separation of an agent's identity, its core instructions, and the tools it can access facilitates the creation of highly specialized and intelligent agents capable of addressing a wide range of complex tasks.

### 3.2 Workflow Agents:

Workflow agents in the ADK are specifically designed to facilitate the creation of structured interactions, where the execution of tasks or the interactions between agents follow a predefined sequence or a specific pattern.<sup>4</sup> This type of agent is particularly valuable for automating processes that have a predictable and repeatable flow of actions.<sup>4</sup> The ADK provides several built-in types of workflow agents to handle different common patterns of structured interactions:

- 4: Sequential Agents are designed to execute a series of steps or interact with other agents in a linear, predefined order.<sup>4</sup> In this type of workflow, each step must be completed before the subsequent step can begin, creating a predictable and straightforward pipeline of actions.<sup>4</sup>
- Loop Agents introduce the capability of repetition. These agents allow for the execution of certain steps or interactions multiple times based on specific conditions that are evaluated during the workflow.<sup>4</sup> This is particularly useful for tasks that require iterative processing, such as refining a result through repeated steps, or for continuous monitoring of a condition until a certain criterion is met.<sup>4</sup>
- Parallel Agents enable the simultaneous execution of multiple independent tasks or interactions.<sup>4</sup> This can significantly reduce the overall processing time for complex tasks that can be broken down into sub-tasks that do not depend on each other's completion.<sup>4</sup>

Workflow agents are especially useful in scenarios where a consistent and repeatable sequence of actions is required.<sup>4</sup> Examples of such scenarios include automated customer service flows that follow a specific script, data processing pipelines that involve a series of transformations, or business process automation where a predefined set of steps needs to be executed.<sup>4</sup> The inclusion of specific workflow agent types like Sequential, Loop, and Parallel showcases the ADK's ability to handle structured and deterministic processes. This capability complements the more dynamic and reasoning-based nature of LLM agents, providing developers with a comprehensive toolkit for building a wide range of AI applications. By offering these pre-defined workflow agent types, the ADK provides developers with ready-made solutions for common structured interaction patterns, simplifying the development of applications that require predictable execution flows and reducing the need to implement these patterns from scratch.

### 3.3 Custom Agents:

Recognizing that the diverse landscape of AI agent applications will inevitably include scenarios that do not perfectly align with predefined agent types, the ADK offers the essential



flexibility to build custom agents with unique functionalities specifically tailored to meet highly specific requirements.<sup>7</sup> The ADK's fundamental "code-first" approach inherently supports the creation of custom agents by empowering developers to define their own agent logic and behavior directly through the use of Python code.<sup>3</sup> This direct control over the agent's implementation provides a very high degree of customization, allowing developers to create agents that precisely match their intended functionality.<sup>3</sup> Developers can further extend the capabilities of their custom agents by integrating their own user-defined tools.<sup>3</sup> This allows the agents to interact with specific external APIs or perform specialized actions that are not covered by the standard set of built-in or readily available third-party tools.<sup>3</sup> Custom agents also offer the ability to implement unique memory management strategies that can be specifically designed to suit the particular needs of the application.<sup>3</sup> This goes beyond the standard memory options provided by the ADK, allowing for more fine-grained control over how the agent stores and retrieves information. Examples of scenarios that might necessitate the development of custom agents include applications in highly specialized domains with unique requirements, agents that need to follow interaction models that deviate from standard conversational patterns, or agents that need to integrate with proprietary or legacy systems that have specific communication protocols. The ability to create custom agents underscores the ADK's commitment to providing a flexible and adaptable platform for AI agent development. This capability allows developers to address a wide range of unique application requirements that might not be met by the standard agent types. The combination of the "code-first" approach and the ability to integrate custom tools and memory management strategies provides the necessary flexibility for even the most specialized agent applications.

### 3.4 Multi-Agent Systems:

A central and defining characteristic of the Google Agent Development Kit is its strong emphasis on enabling the development and seamless coordination of multiple interacting AI agents.<sup>7</sup> This focus is driven by the understanding that many complex real-world problems are beyond the capabilities of a single, monolithic agent and require the collaborative efforts of multiple specialized entities.<sup>7</sup> Building applications that are both modular and highly scalable through the composition of multiple specialized agents organized within a hierarchical structure is a core design principle of the ADK.<sup>1</sup> This architectural approach allows for the creation of sophisticated systems where different agents are assigned specific responsibilities based on their expertise, leading to better overall system organization, improved maintainability of the codebase, and enhanced reusability of individual agent components across different parts of the application or even in other projects.<sup>1</sup> The ADK provides robust mechanisms for enabling intricate coordination and delegation of tasks between these specialized agents.<sup>1</sup> For instance, a primary agent within a multi-agent system can analyze a user's request and then intelligently delegate specific sub-tasks to other agents in the system based on their described skills and areas of expertise.<sup>1</sup> The framework facilitates intelligent routing of user requests to the most appropriate agent within the multi-agent system.<sup>1</sup> This routing decision is typically based on a careful analysis of the content of the user's request and a comparison with the documented capabilities of each available agent in the system, ensuring that each task is handled by the agent best suited to

perform it efficiently and effectively.<sup>1</sup> The ADK also streamlines the process of transferring control and planning between different agents within a multi-agent system.<sup>6</sup> This seamless transition of responsibility is particularly crucial for complex workflows that require a series of interconnected steps or involve different areas of expertise at various stages of the process.<sup>6</sup> Remarkably, the ADK allows developers to build basic yet functional multi-agent systems with a relatively small amount of code, often under 100 lines of Python.<sup>6</sup> This demonstrates the framework's efficiency and ease of use when it comes to creating collaborative AI applications, making it accessible even for developers who are new to the concept of multi-agent systems.<sup>6</sup> The tutorial provided within the ADK documentation offers a practical, step-by-step guide that walks developers through the entire process of building a multi-agent system.<sup>7</sup> This hands-on approach further illustrates the framework's capabilities in this area and provides developers with a solid foundation for building their own multi-agent applications.<sup>7</sup> The ADK truly shines when it is utilized to construct collaborative multi-agent systems that can effectively leverage a diverse range of tools to accomplish complex objectives.<sup>1</sup> The ability of multiple agents to work together in a coordinated manner and to utilize various resources, both internal and external, is a key strength of the framework and enables the development of highly sophisticated AI solutions.<sup>1</sup> The framework supports the creation of hierarchical structures for organizing agents within a multi-agent system.<sup>1</sup> This allows for the definition of parent-child relationships between agents, where a primary agent can take on the role of managing and delegating tasks to a set of sub-agents.<sup>1</sup> This hierarchical approach is particularly beneficial for managing the inherent complexity of large-scale multi-agent systems, providing a clear and organized way to structure the interactions and responsibilities of the various agents involved.<sup>1</sup> Consider a practical example within a customer service application: a primary agent could be responsible for initially interacting with the user and understanding their needs. Based on the user's request, this primary agent could then delegate specific tasks, such as verifying the user's identity, retrieving their order information from a database, and processing a return request, to specialized sub-agents designed for each of these specific functions.<sup>1</sup> The ADK makes it relatively straightforward to define these relationships between agents and to manage the flow of information and control between them, ensuring a cohesive and efficient process for handling the user's request.<sup>1</sup> Complementing the ADK is the introduction of the Agent2Agent (A2A) protocol, a new open standard that further enhances the capabilities of multi-agent systems.<sup>3</sup> The A2A protocol enables AI agents to communicate with each other securely, exchange vital information, and coordinate their actions across a wide range of enterprise platforms and applications, regardless of the specific underlying technology stack or tools that were used to build them.<sup>3</sup> This focus on interoperability promises to unlock a new era of collaboration between AI agents, fostering innovation and leading to the development of even more powerful and versatile multi-agent systems capable of tackling increasingly complex challenges.<sup>3</sup> The profound emphasis on multi-agent systems within the ADK, coupled with features like the ability to create hierarchical agent structures, the implementation of intelligent request routing, and the introduction of the groundbreaking A2A protocol, firmly establishes the ADK as a leading framework for developing highly sophisticated and collaborative AI applications. This strategic direction underscores the belief that the future of

advanced AI lies in the synergistic interactions of multiple intelligent agents working together. The ADK is clearly designed to be a foundational toolkit for realizing this vision, providing developers with the necessary tools and architectural patterns to build the next generation of intelligent and collaborative AI systems.

## **4. Integrating Language Models:**

### **4.1 Compatible Models:**

The Google ADK is intentionally designed to be a versatile framework that can seamlessly integrate and operate with a wide variety of language models.<sup>4</sup> This flexibility allows developers to select the most appropriate language model based on the specific requirements and characteristics of their AI agent application.<sup>4</sup> The framework is particularly optimized for Google's own family of advanced Gemini models, which represent the state-of-the-art in various artificial intelligence tasks.<sup>1</sup> This deep level of optimization ensures that developers can fully leverage the cutting-edge capabilities offered by the Gemini models within their ADK-based agents, potentially achieving superior performance and access to the latest advancements in Google's AI technology.<sup>1</sup> The ADK features built-in integration with Vertex AI Model Garden, a comprehensive platform that provides access to an extensive selection of pre-trained models not only from Google but also from other leading artificial intelligence providers.<sup>1</sup> This integration significantly simplifies the process for developers to deploy and subsequently utilize a diverse range of models within their ADK agents, as Vertex AI offers a managed environment for hosting, scaling, and managing these models.<sup>1</sup> To further expand the range of language models that can be used with the ADK, the framework incorporates seamless integration with LiteLLM.<sup>1</sup> LiteLLM is an open-source library that provides a unified and consistent interface for interacting with models from various prominent providers, including Anthropic, Meta, Mistral AI, and AI21 Labs.<sup>1</sup> This integration allows developers to easily experiment with and utilize models from different ecosystems within their ADK projects without having to manage the intricacies of each provider's specific API.<sup>1</sup> In general, the ADK is designed to be compatible with a wide array of popular Large Language Models (LLMs) and open-source generative AI tools, providing developers with a broad spectrum of options to choose from based on their specific needs and preferences.<sup>4</sup> The ADK's overall design philosophy emphasizes tight integration with the Google ecosystem, particularly with Gemini models, making it an especially compelling choice for developers who are looking to leverage Google's advanced artificial intelligence capabilities.<sup>4</sup> The framework is engineered to make it straightforward for developers to get started with building simple agents that are powered by Gemini models and other Google AI tools, while simultaneously providing the necessary control and architectural structure required for developing more complex agent architectures that might utilize other language models.<sup>4</sup> The quickstart guide often includes practical examples that demonstrate how to use specific models, such as `gemini-2.0-flash-exp`, to illustrate the process of configuring and interacting with language models within the ADK framework.<sup>5</sup> Similarly, the tutorial that guides developers through the creation of multi-agent systems also showcases the use of specific models through the LiteLLM integration, for instance, `MODEL_GPT_4O`, further highlighting the framework's inherent flexibility in supporting different models for various agents within a single

application.<sup>8</sup> The ADK's extensive compatibility with a diverse range of language models, achieved through its native support for Gemini and Vertex AI Model Garden, as well as its seamless integration with LiteLLM, provides developers with an unparalleled level of flexibility in selecting the optimal model for their specific AI agent applications. This multi-faceted approach to model integration ensures that developers are not locked into a single provider or model type, but rather have the freedom to choose the best tool for the job based on factors such as performance, cost, and specific task requirements.

#### 4.2 Integration Methods:

The ADK offers direct integration with Gemini models through its own API, allowing developers to easily access and utilize the advanced capabilities of these models within their agents. This native integration often provides optimized performance and access to the latest Gemini features. By integrating with Vertex AI Model Garden, the ADK allows developers to seamlessly deploy and access a wide variety of models from Google and other providers. Vertex AI provides a managed platform for model hosting, scaling, and management, simplifying the process of using different models in ADK agents. The integration with LiteLLM enables ADK developers to connect to models from various providers (like OpenAI, Anthropic, Mistral) through a consistent and unified API. This abstraction layer simplifies the process of working with different model APIs and allows for easy switching between models if needed. Within the ADK, developers typically configure the desired language model and its associated API keys or credentials directly within the definition of their agent. This allows for fine-grained control over which model each agent in a multi-agent system utilizes. This multi-faceted approach to model integration provides developers with the flexibility to choose the most suitable language model based on factors such as reasoning capabilities, cost, latency requirements, and specific task performance. The ADK's architecture supports leveraging different models for different agents or even switching models during the development process to optimize performance and cost.

### 5. Harnessing the Power of Tools:

#### 5.1 Function Tools:

The ADK provides developers with the capability to significantly extend the functionality of their AI agents by integrating custom Python functions as tools.<sup>3</sup> This powerful feature enables agents to perform specific actions or interact with external systems by simply calling these user-defined functions.<sup>3</sup> Function tools are particularly valuable in scenarios where agents need to interact with external APIs to retrieve or manipulate data, perform specialized data processing tasks that are not inherently part of the language model's capabilities, or execute any other custom logic that is specific to the application's requirements.<sup>3</sup> The tutorial included in the ADK documentation provides a clear demonstration of how to define simple Python functions, such as `say_hello` for generating greetings and `say_goodbye` for farewell messages, and then register these functions as tools that an agent can utilize during its operation.<sup>8</sup> When defining function tools within the ADK, it is crucial to provide clear and concise docstrings that accurately describe the tool's purpose, the arguments it accepts as input, and the type of value it returns as output.<sup>8</sup> This documentation is essential as it helps the agent understand how and when to effectively use the tool to accomplish its tasks.<sup>8</sup>

## 5.2 Built-in Tools:

The ADK comes with a set of pre-packaged tools that offer agents immediate access to common and essential functionalities right from the start.<sup>7</sup> These built-in tools are designed to provide core capabilities without requiring developers to implement them from scratch.<sup>7</sup> Examples of such built-in tools include a tool for performing web searches, such as `google_search`, which allows agents to retrieve information from the internet, and a tool for executing code snippets, enabling agents to run dynamic code and potentially interact with their environment.<sup>2</sup> The quickstart guide for the ADK often demonstrates the practical use of these built-in tools through examples, such as tools for retrieving the current weather (`get_weather`) and the current time (`get_current_time`) for a specified location.<sup>11</sup> These examples clearly illustrate how easily agents can be equipped with basic yet highly useful capabilities that enhance their ability to respond to user queries and perform tasks.<sup>11</sup>

## 5.3 Third-Party Tools:

Recognizing the vast and ever-growing ecosystem of existing tools, libraries, and services available to developers, the ADK allows for the seamless integration of various third-party tools and services to further expand the functionality of AI agents built with the framework.<sup>7</sup> This capability is crucial for leveraging specialized tools that might not be available as built-in features of the ADK.<sup>7</sup> The framework offers seamless integration with popular AI development frameworks such as LangChain and CrewAI.<sup>4</sup> This integration allows developers to readily utilize the extensive collections of tools and utilities provided by these libraries within their ADK-based agents, significantly broadening the range of tasks that the agents can perform.<sup>4</sup> Developers can also easily integrate with libraries like LlamaIndex, which specializes in knowledge retrieval and integration from diverse data sources.<sup>2</sup> This enables ADK agents to access and process information from a wide variety of data formats and locations, enhancing their ability to answer complex questions and perform information-intensive tasks.<sup>2</sup> The process of integrating third-party tools typically involves importing the necessary libraries into the project and then configuring the agent to use the desired tools by referencing them in the agent's definition.<sup>2</sup> The ADK provides the underlying mechanisms to handle the interaction between the agent and these external tools, ensuring a smooth and efficient integration process.<sup>2</sup>

## 5.4 Google Cloud Tools:

Given its origin and close relationship with the Google Cloud platform, the ADK provides strong and seamless integration with various services offered within Google Cloud.<sup>7</sup> This tight integration allows agents built with the ADK to leverage the immense power and scalability of Google's cloud infrastructure.<sup>7</sup> A key integration is with Vertex AI, which not only provides access to a wide range of cutting-edge language models but also offers comprehensive services for deploying, managing, and continuously monitoring ADK agents in a production environment.<sup>1</sup> The ADK enables agents to directly connect with APIs and connectors that are managed within the Google Cloud platform.<sup>6</sup> This allows agents to interact with a diverse set of Google Cloud services and access various data sources hosted within the cloud, facilitating the development of cloud-native AI applications.<sup>6</sup> The framework is specifically optimized for seamless integration within the broader Google Cloud ecosystem.<sup>1</sup> This makes the ADK a natural and highly advantageous choice for developers and organizations that are already

building or planning to build their AI applications on the Google Cloud platform, as it provides a cohesive and efficient development experience.<sup>1</sup>

#### 5.5 MCP Tools:

The ADK incorporates support for the Model Context Protocol (MCP), which is an open standard designed to facilitate the connection of AI agents to a vast and diverse range of data sources and existing capabilities.<sup>2</sup> By including support for MCP, the ADK allows agents to seamlessly leverage a growing ecosystem of tools that are compatible with the MCP standard.<sup>6</sup> This enables agents to access and interact with a wide variety of data and services in a standardized and efficient manner, regardless of the underlying implementation details.<sup>6</sup> MCP is particularly beneficial for enabling agents to efficiently reason over extremely large datasets.<sup>6</sup> It achieves this by allowing agents to access and process information through storage artifacts, rather than relying solely on the potentially limited context window of the underlying language model.<sup>6</sup> This approach can significantly improve the agent's ability to handle complex queries and tasks that require processing large amounts of information.<sup>6</sup> The ADK's support for MCP ensures seamless connectivity to tools from a multitude of diverse sources, further expanding the already extensive range of capabilities that can be integrated into AI agents built with the framework.<sup>6</sup>

#### 5.6 OpenAPI Tools:

The ADK provides robust mechanisms for integrating with external services that expose their functionality through the OpenAPI specification (formerly known as Swagger).<sup>5</sup> OpenAPI is a widely adopted standard for describing and documenting RESTful APIs, making it easier for different systems to understand and interact with each other.<sup>5</sup> By leveraging OpenAPI specifications, the ADK can automatically understand the capabilities and interfaces of external services.<sup>5</sup> This allows agents to automatically generate the necessary API calls and effectively handle the responses they receive from these services, without requiring developers to manually write the integration code for each specific API endpoint.<sup>5</sup> The ability to seamlessly integrate with any service that provides an OpenAPI endpoint significantly expands the potential capabilities of ADK agents.<sup>6</sup> It allows them to interact with a vast array of external systems and data sources across the internet in a standardized and efficient manner, opening up a wide range of possibilities for building sophisticated and interconnected AI applications.<sup>6</sup> The comprehensive support within the ADK for various types of tools and integration methods empowers developers to equip their AI agents with an extensive array of capabilities. These range from executing custom Python code and performing real-time web searches to seamlessly interacting with a multitude of third-party services, leveraging the power of the Google Cloud infrastructure, and connecting with systems that adhere to open standards like MCP and OpenAPI. This rich and diverse tooling ecosystem is undoubtedly a key strength of the ADK, providing developers with the flexibility and resources they need to build highly functional and versatile AI applications.

## **6. Securing Your Agents: Authentication Strategies:**

#### 6.1 Authentication Mechanisms:

Security is of paramount importance in the development of AI agents, particularly those that handle sensitive data or perform critical actions.<sup>4</sup> The ADK recognizes this need and provides

several robust mechanisms to ensure the security of agent interactions and to carefully control access to various resources.<sup>4</sup> The quickstart guide for the ADK emphasizes the crucial step of setting up appropriate credentials for accessing Google Cloud services.<sup>10</sup> This typically involves using the Google Cloud CLI command `gcloud auth application-default login` to generate local Application Default Credentials (ADC).<sup>10</sup> These credentials are then used by the agent to authenticate with Vertex AI during the local development phase, ensuring secure access to the necessary cloud resources.<sup>10</sup>

## 6.2 Agent Auth:

Agent authentication is the process of defining the specific identity under which an AI agent operates within a system.<sup>17</sup> This identity is crucial as it determines the permissions and the range of resources that the agent is authorized to access and utilize.<sup>17</sup> By carefully controlling and defining the agent's identity, developers can establish strict security boundaries, ensuring that agents can only perform actions that they have been explicitly allowed to carry out.<sup>17</sup> This principle of least privilege is fundamental to building secure and reliable AI applications.<sup>17</sup>

## 6.3 User Auth:

User authentication, in the context of the ADK, focuses on the critical process of verifying the identity of the user who is currently interacting with the AI agent.<sup>17</sup> This is particularly important in scenarios where the agent needs to perform actions on behalf of a specific user or access data that is specific to that user's account or profile.<sup>17</sup> In the ADK, user authentication is often implemented using the industry-standard OAuth 2.0 protocol.<sup>17</sup> This protocol allows the agent to securely interact with a frontend application (such as a web or mobile interface) to acquire an OAuth token.<sup>17</sup> This token serves as a digital proof of the user's identity and their authorization to perform certain actions when the agent attempts to interact with external services or access protected resources.<sup>17</sup> A significant advantage of employing user authentication is that agents will only perform actions that the authenticated user themselves would have the permission to perform.<sup>17</sup> This greatly reduces the potential risk of a malicious user being able to exploit the agent to gain unauthorized access to additional data or functionalities beyond the scope of their own authorized permissions.<sup>17</sup>

## 6.4 Other Security Considerations:

In addition to robust authentication mechanisms, the ADK incorporates several other important security measures to protect against potential vulnerabilities and ensure the integrity of AI agent applications <sup>17</sup>:  
**Guardrails:** The framework allows developers to implement guardrails, which are mechanisms to carefully screen both the inputs received by the agent from users and the outputs that the agent generates in response.<sup>17</sup> This helps to control the behavior of the underlying language model and any tools that the agent utilizes, ensuring that they operate within predefined boundaries and adhere to established security policies.<sup>17</sup>  
**In-Tool Guardrails:** Developers have the ability to design tools defensively by using developer-set tool context to enforce specific policies.<sup>17</sup> For example, a tool designed to query a database could be configured to only allow queries on a specific set of tables, thereby limiting the agent's ability to access potentially sensitive information stored in other parts of the database.<sup>17</sup>  
**Gemini Safety Features:** When developers choose to use Gemini models within the ADK, they can leverage the built-in content safety filters provided by

Google.<sup>17</sup> These filters are designed to automatically block the output of content that is deemed harmful, inappropriate, or violates safety guidelines.<sup>17</sup> Additionally, developers can utilize system instructions to directly guide the model's behavior and proactively enforce safety guidelines that are specific to their application's needs.<sup>17</sup> Model and Tool Callbacks: The ADK provides the flexibility to define and implement callbacks, which are custom functions that can be executed before or after the invocation of a language model or the execution of a tool.<sup>17</sup> These callbacks can be used to perform various security-related tasks, such as validating the parameters of a function call or inspecting the content of a model's response against predefined criteria or external policies, providing an additional layer of security and control.<sup>17</sup> Sandboxed Code Execution: To mitigate the potential security risks associated with allowing a language model to generate and execute code, the ADK supports sandboxed code execution environments.<sup>17</sup> This isolates the execution of any generated code within a restricted environment, limiting its ability to interact with the underlying system and preventing it from causing unintended harm or security breaches.<sup>17</sup> Evaluation and Tracing: The ADK's built-in evaluation tools can be used to assess not only the functional correctness but also the security implications of an agent's output.<sup>17</sup> Furthermore, tracing capabilities provide detailed visibility into the sequence of actions taken by an agent during its execution, including its choice of tools and its reasoning process.<sup>17</sup> This level of transparency is invaluable for identifying and analyzing potential security vulnerabilities or unexpected behaviors.<sup>17</sup> Network Controls and VPC-SC: For deployments in the Google Cloud environment, the ADK can be used in conjunction with network controls and VPC Service Controls (VPC-SC).<sup>17</sup> These security features allow developers to confine the agent's activity within secure network perimeters, preventing unauthorized access to resources and limiting the potential impact of any security breaches or data exfiltration attempts.<sup>17</sup> The ADK's provision of a comprehensive suite of authentication and security mechanisms, including agent and user authentication, guardrails, sandboxed execution, and integration with Gemini's safety features, clearly demonstrates a strong commitment to enabling developers to build secure and trustworthy AI agent applications. This multi-faceted approach to security, addressing various aspects from identity management to runtime protection, highlights the framework's maturity and its consideration for the critical security implications of deploying AI agents in real-world scenarios.

## **7. Deploying ADK Agents:**

### **7.1 Agent Engine:**

Agent Engine is a fully managed runtime environment offered within Google Cloud's Vertex AI platform, specifically designed to streamline the deployment of custom-built AI agents, including those developed using the ADK, into production environments.<sup>4</sup> This managed service significantly reduces the operational overhead associated with deploying and running AI agents at scale by handling many of the underlying complexities.<sup>6</sup> These complexities include managing the agent's conversational context, automatically provisioning and managing the necessary infrastructure resources, dynamically scaling the agent's capacity to handle fluctuating workloads, ensuring robust security measures are in place, providing integrated tools for evaluating agent performance, and offering comprehensive monitoring



capabilities to track the agent's health and behavior in real-time.<sup>6</sup> Agent Engine is engineered to seamlessly integrate with the ADK, as well as with other popular agent development frameworks, providing a smooth and efficient "develop-to-deploy" experience.<sup>6</sup> This tight integration minimizes the friction typically encountered when transitioning AI agents from the development phase to a fully operational production system.<sup>6</sup> By leveraging Agent Engine, developers can benefit from a fully managed, highly scalable, and enterprise-grade runtime environment.<sup>1</sup> This ensures that their AI agents can reliably and securely handle the demands of production-level workloads, meeting the performance and availability expectations of enterprise applications.<sup>1</sup> Agent Engine offers built-in features that further simplify the deployment process, including capabilities for testing agent deployments in a staging environment, managing the release of new agent versions, and ensuring high levels of reliability and availability at a global scale.<sup>6</sup> These features are crucial for organizations that need to deploy and maintain their AI agents with confidence and minimal operational burden.<sup>6</sup> Agent Engine operates within the context of Google's AgentSpace, a platform designed to provide organizations with a centralized hub for discovering, managing, and accessing a variety of AI agents.<sup>9</sup> This integration further enhances the overall experience of deploying and utilizing AI agents within the Google Cloud ecosystem.<sup>9</sup>

## 7.2 Cloud Run:

Google Cloud Run provides another powerful and flexible option for deploying AI agents developed using the ADK.<sup>4</sup> It is a serverless container platform that allows developers to run stateless containers on a fully managed environment, abstracting away the complexities of server management.<sup>4</sup> To deploy an ADK agent using Cloud Run, developers typically containerize their agent application using Docker.<sup>4</sup> This involves creating a Dockerfile, which is a script that specifies the exact environment and all the necessary dependencies required to run the agent application.<sup>4</sup> Once the Dockerfile is defined, a container image is built based on this specification. This container image can then be easily deployed to Google Cloud Run.<sup>4</sup> Cloud Run offers several key benefits that make it an attractive option for deploying AI agents.<sup>5</sup> These benefits include automatic scaling of the application based on the volume of incoming requests, a pay-as-you-go pricing model where you only pay for the resources consumed by your running containers, and a fully managed infrastructure that eliminates the need for developers to provision and manage underlying servers.<sup>5</sup> This serverless approach allows developers to focus solely on building and deploying their AI agent logic without being burdened by infrastructure management tasks.<sup>5</sup>

## 7.3 Local Deployment and Custom Infrastructure:

For the purposes of initial development, thorough testing, and rapid prototyping, ADK agents can be conveniently run locally on a developer's own machine.<sup>4</sup> This local deployment capability allows for quick iteration cycles and efficient debugging in a controlled and isolated environment before the agent is deployed to a cloud platform.<sup>4</sup> Organizations may also have specific requirements or existing infrastructure that necessitate deploying ADK agents within their own custom environments, which could include on-premises servers or other cloud platforms.<sup>1</sup> The ADK's inherent support for containerization using Docker makes it relatively straightforward to package and deploy agents in a wide variety of environments.<sup>1</sup> By creating a Docker image of the ADK agent application, organizations can ensure consistency and

portability across different deployment targets, whether it's their own private infrastructure or a different cloud provider.<sup>1</sup> The availability of a range of flexible deployment options within the ADK caters to different needs and stages of development. From the fully managed Agent Engine for production scalability and the container-based flexibility of Cloud Run to the convenience of local execution for development and testing, the ADK provides developers with the tools to deploy their AI agents in the way that best suits their specific requirements and infrastructure.

## **8. Managing Conversations and Agent State:**

### **8.1 Session Management:**

Session management within the ADK refers to the mechanisms by which the framework handles individual user interactions with an AI agent.<sup>5</sup> A key aspect of session management is the ability to maintain context across multiple turns within a single conversation.<sup>5</sup> This allows the agent to remember previous parts of the dialogue, such as user preferences or the topic of discussion, and to respond to subsequent inputs in a coherent and contextually relevant manner.<sup>5</sup> Effective session management is crucial for creating engaging and natural conversational experiences for users.<sup>5</sup>

### **8.2 State Management:**

State management in the ADK deals with how the framework enables AI agents to maintain their internal state across different interactions or user sessions.<sup>2</sup> This capability is essential for building agents that can remember information, track their progress through a multi-step task, or adapt their behavior based on past events or user preferences.<sup>2</sup> The ADK provides developers with the tools and mechanisms necessary for storing and retrieving the agent's internal state.<sup>2</sup> This allows important information to persist beyond the scope of a single user interaction, ensuring that the agent can retain knowledge and context over time.<sup>2</sup> This state can include various types of data, such as user preferences that were explicitly stated in previous interactions, the current stage of a complex task that the user is trying to accomplish with the agent, or any other relevant information that the agent needs to retain to provide a consistent and personalized experience.<sup>2</sup>

### **8.3 Memory:**

The ADK offers different types of memory that AI agents can utilize to store and subsequently recall information as needed.<sup>7</sup> These memory capabilities are crucial for enabling agents to maintain context and learn from past interactions.<sup>7</sup> Short-term memory is typically employed to maintain the immediate context of an ongoing conversation.<sup>3</sup> It allows the agent to remember the recent turns of dialogue within a single session, enabling it to understand references to previous statements and maintain coherence throughout the conversation.<sup>3</sup> Long-term memory, on the other hand, enables agents to store and recall information over extended periods of time or across multiple distinct user sessions.<sup>3</sup> This type of memory can be used to remember user preferences that were established in previous interactions, recall details from past conversations that might be relevant to the current one, or retain knowledge that the agent has acquired over time through its interactions and experiences.<sup>3</sup> The ADK provides developers with the tools and configurations necessary to implement and manage both short-term and long-term memory for their agents.<sup>3</sup> Developers have the flexibility to

choose and implement different memory strategies depending on the specific requirements of their agent application.<sup>3</sup> This involves balancing the need for the agent to have a strong sense of conversational context with the efficiency and scalability of the underlying memory storage mechanisms.<sup>3</sup>

#### 8.4 Artifacts:

The concept of artifacts within the ADK appears to be related to the management and handling of various forms of data and files that are associated with the interactions between users and AI agents.<sup>7</sup> Interestingly, under the "Artifacts" category in the documentation outline, the entry is "None".<sup>4</sup> This might suggest that this specific feature is either not yet fully implemented within the ADK, or that the documentation pertaining to artifacts is currently limited or under development.<sup>4</sup> However, one of the research snippets mentions that the ADK enables agents to efficiently process and reason over very large datasets through the utilization of storage artifacts.<sup>6</sup> This suggests that there is indeed a mechanism within the ADK for handling and processing external data in a way that is optimized for agent performance.<sup>6</sup> Additionally, the ADK is mentioned as being capable of handling file uploads and downloads.<sup>5</sup> These capabilities could also be considered a form of artifact management, as they involve the transfer and handling of external files as part of the agent's interactions with users or other systems.<sup>5</sup>

#### 8.5 Callbacks:

While callbacks in the ADK serve primarily as a mechanism for extending the core functionality of AI agents (as will be discussed in more detail in the subsequent section), they can also play a significant role in the management of agent state and memory.<sup>4</sup> Callbacks allow developers to execute custom code at various critical points in the agent's lifecycle, such as before or after the agent processes a user input or generates a response.<sup>4</sup> By strategically using callbacks, developers can implement custom logic for updating the agent's internal state or for storing specific pieces of information in the agent's memory based on the particular events that occur during an interaction.<sup>4</sup> This provides a flexible way to manage the agent's context and ensure that it retains the necessary information to function effectively.<sup>4</sup> The ADK provides a foundational framework for effectively managing the flow of conversations through sessions, maintaining the agent's state for persistence across interactions, and leveraging different types of memory to enhance the agent's intelligence and ability to provide contextually relevant responses. While the specific role and full functionality of artifacts within the ADK appear to be an area that might see further development and clearer documentation in the future, the overall framework provides developers with the essential tools to manage the context and persistence of their AI agent applications.

### **9. Extending Functionality with Callbacks:**

#### 9.1 Types of Callbacks:

Callbacks in the ADK offer a flexible and powerful way for developers to intercept and extend the default behavior of AI agents at various key stages of their execution.<sup>4</sup> By defining and registering custom callback functions, developers can effectively inject their own specific logic into the agent's processing pipeline, allowing for a high degree of customization.<sup>4</sup> The ADK likely provides different types of callbacks that are triggered at specific points in the

agent's lifecycle.<sup>4</sup> These could include callbacks that are executed immediately before a language model is invoked to generate a response, callbacks that are triggered right after the model has produced its output, callbacks that are executed just before a tool is called by the agent, or callbacks that are invoked after a tool has completed its execution and returned its result.<sup>4</sup> The precise types of callbacks that are available within the ADK and the specific points in the execution flow at which they are triggered would be detailed in the framework's comprehensive API reference.<sup>4</sup>

## 9.2 Callback Patterns:

The ADK documentation mentions the existence of various callback patterns that developers can follow to implement common and useful extensions to the standard behavior of their AI agents.<sup>4</sup> These patterns likely provide guidance and best practices on how to structure callback functions effectively to achieve specific desired outcomes and functionalities.<sup>4</sup> One particularly important use case for callbacks within the ADK is to perform validation of both the inputs and the outputs of the language model and any tools that the agent utilizes during its operation.<sup>17</sup> By implementing callback functions that check the parameters of a function call before it is executed or that examine the content of a model's response after it has been generated, developers can ensure that the agent behaves as expected and adheres to certain predefined constraints or policies.<sup>17</sup> Callbacks can also be effectively used to implement safety guardrails within AI agents.<sup>17</sup> For instance, a callback function could be defined to screen the input provided by a user to the agent or to filter the output generated by the language model, checking for any potentially harmful, inappropriate, or undesirable content.<sup>17</sup> This allows developers to build more responsible and trustworthy AI agents that are less likely to produce problematic responses or take unintended actions.<sup>17</sup> The "Responsible Agents" guide within the ADK documentation specifically suggests using callbacks in conjunction with a cheap and fast language model, such as Gemini Flash Lite, to implement an additional layer of safety by screening both the inputs received by the agent and the outputs it generates.<sup>17</sup> This pattern provides an efficient and cost-effective way to perform content moderation and ensure the safety of the agent's interactions.<sup>17</sup> The use of callbacks in the ADK offers a highly flexible and powerful approach to deeply customize the behavior of AI agents. By allowing developers to inject their own custom logic at various strategic stages of the agent's execution flow, including crucial steps like input and output validation and safety checks, callbacks enable the creation of robust, secure, and highly tailored AI applications that can meet a wide range of specific requirements.

## 10. Understanding the ADK Runtime Environment:

### 10.1 Events:

The ADK's runtime environment likely incorporates a robust mechanism for handling events that occur during the execution of an AI agent.<sup>4</sup> These events could represent a variety of occurrences within the agent's operational lifecycle, such as the successful invocation of an external tool, the generation of a response by the underlying language model, or the successful completion of a specific sub-task within a larger workflow.<sup>4</sup> Having a clear understanding of how events are handled within the ADK is crucial for developers who want to build agents that can react dynamically to different situations and effectively manage their

own execution flow.<sup>4</sup> The official ADK documentation would likely provide detailed information on the specific types of events that are emitted by the framework during agent execution and how developers can subscribe to and handle these events using appropriate event handling mechanisms.<sup>4</sup>

## 10.2 Context:

The concept of context is absolutely fundamental to the effective operation of AI agents.<sup>4</sup> In the ADK, context refers to all the relevant information that is readily available to the agent at any given point during its execution.<sup>4</sup> This crucial information can encompass a wide range of data, including the user's current input or query, the complete history of the ongoing conversation between the user and the agent, the agent's current internal state, which might include variables or data it has stored, and any relevant environmental variables or parameters that might influence the agent's behavior.<sup>4</sup> Agent Engine, the managed deployment service provided by Google Cloud for ADK agents, plays a significant role in handling and effectively managing the context of agents that are deployed using this service.<sup>6</sup> This ensures that the deployed agents have continuous access to the necessary information they need to operate correctly and provide contextually appropriate responses to user interactions.<sup>6</sup> Furthermore, the ADK allows developers to set tool-specific context.<sup>17</sup> This feature enables developers to provide specific information or enforce certain policies that are relevant only when a particular tool is invoked by the agent.<sup>17</sup> For example, when an agent uses a database query tool, the tool-specific context might include restrictions on which tables the agent is allowed to access or specific parameters that must be included in the query.<sup>17</sup> This capability provides a fine-grained level of control over how tools are used and helps ensure that they operate within defined security and functional boundaries.<sup>17</sup> The ADK's runtime environment, with its sophisticated handling of events and its robust management of context, provides the essential infrastructure for building dynamic and context-aware AI agents. These agents can effectively respond to user inputs by leveraging the relevant information at hand and manage their internal state appropriately throughout their execution.

## 11. Evaluating and Improving Your Agents:

### 11.1 Evaluation Methods:

The ADK places a significant emphasis on the ability for developers to systematically evaluate the performance and overall effectiveness of their AI agents.<sup>4</sup> This focus on evaluation is absolutely crucial for ensuring that the agents are functioning correctly, providing accurate and relevant responses to user queries, and ultimately meeting the specific objectives for which they were designed.<sup>4</sup> The evaluation process within the ADK is designed to be comprehensive, involving the assessment of both the quality of the final response generated by the agent and the detailed, step-by-step execution trajectory that the agent followed to arrive at that particular response.<sup>4</sup> This dual approach provides developers with a holistic view of the agent's performance, allowing them to understand not only the outcome but also the reasoning and the process behind it.<sup>4</sup> Developers can rigorously evaluate their agents by testing them against a set of predefined test cases.<sup>4</sup> These test cases typically consist of specific input prompts or scenarios along with the corresponding expected outputs or desired behaviors.<sup>4</sup> By running these test cases and comparing the agent's actual responses

and actions to the expected outcomes, developers can obtain objective measurements of the agent's accuracy, reliability, and overall performance across a range of different situations.<sup>4</sup>

#### 11.2 Evaluation Tools:

The ADK provides a suite of built-in tools and a structured framework to facilitate the entire evaluation process for developers.<sup>1</sup> These tools likely enable developers to easily define their test cases, efficiently run the evaluations against their AI agents, and then thoroughly analyze the resulting data to identify specific areas where the agent's performance might need improvement or refinement.<sup>1</sup> Agent Engine, the managed deployment service for ADK agents within Google Cloud, also incorporates features for evaluating the performance of agents that have been deployed to a production environment.<sup>6</sup> This allows developers to gain valuable insights into how their agents are behaving in real-world usage scenarios and to identify any potential issues or areas for optimization.<sup>6</sup> Furthermore, the ADK supports tracing, which is an incredibly useful capability that allows developers to gain a highly detailed view into the sequence of actions taken by an agent during its execution.<sup>17</sup> This includes information about the specific tools the agent chose to use at each step, the reasoning behind its decisions, and an overall assessment of the efficiency of its approach to solving a particular problem.<sup>17</sup> Tracing is an invaluable tool for debugging issues that might arise during agent execution and for gaining a deep understanding of exactly how an agent arrives at a particular outcome, which can be crucial for identifying areas where the agent's logic or configuration might need to be adjusted.<sup>17</sup> The ADK documentation likely includes comprehensive guides and detailed explanations on how to effectively utilize these various evaluation methods and tools to thoroughly assess and continuously improve the performance of AI agents built with the framework.<sup>5</sup> The ADK's provision of robust built-in evaluation tools and well-defined methodologies, including the ability to assess both the final responses of agents and their underlying execution trajectories, clearly underscores the framework's strong commitment to helping developers build high-quality and reliable AI agents. This emphasis on systematic evaluation indicates that the ADK is designed with the entire lifecycle of agent development in mind, recognizing that continuous assessment and iterative improvement are fundamental to creating effective and trustworthy AI applications.

### 12. Best Practices and Advanced Guides:

#### 12.1 Responsible Agents:

The ADK documentation includes a dedicated and important guide titled "Responsible Agents," which provides a comprehensive set of best practices, recommendations, and guidelines for building AI agents that are not only powerful and capable but also trustworthy, ethical, and aligned with responsible AI principles.<sup>4</sup> This guide emphasizes the critical importance of proactively implementing responsible AI patterns and carefully considering factors such as fairness, transparency in the agent's decision-making processes, accountability for the agent's actions, and overall safety in the design and development of AI agents.<sup>4</sup> Key topics that are likely covered in the "Responsible Agents" guide include: Identity and Authorization: Ensuring that all agents operate under clearly defined and appropriate identities, and that they possess only the necessary authorizations to perform their intended tasks securely and without exceeding their designated privileges.<sup>17</sup> Guardrails: Implementing



robust mechanisms to control the behavior of agents, preventing them from generating harmful, biased, or otherwise inappropriate content, and ensuring they adhere to predefined ethical and safety boundaries.<sup>17</sup> Sandboxed Code Execution: Utilizing isolated and restricted environments for the execution of any code generated by the agent, thereby protecting against potential security vulnerabilities and preventing unintended or malicious actions.<sup>17</sup> Evaluation and Tracing: Leveraging the ADK's evaluation tools and tracing capabilities to thoroughly understand the agent's behavior, identify potential issues related to bias, fairness, or unintended consequences, and continuously monitor its performance in real-world scenarios.<sup>17</sup> Network Controls and VPC-SC: For agents deployed in cloud environments, utilizing network security measures such as VPC Service Controls (VPC-SC) to confine the agent's activity within secure perimeters, preventing unauthorized access to resources and mitigating the risk of data exfiltration.<sup>17</sup>

## 12.2 Other Guides:

Beyond the crucial guide on responsible agents, the ADK documentation likely contains a collection of other valuable guides that provide detailed explanations of various specific features, advanced functionalities, and best practices for utilizing different aspects of the framework.<sup>4</sup> These additional guides would offer in-depth insights into tasks such as designing effective and natural interactions between users and agents, efficiently managing the state of agents across multiple turns in a conversation, and ensuring the overall safety, reliability, and robustness of agent applications built with the ADK.<sup>8</sup> For developers who are working on more complex applications involving multiple interacting agents, there would likely be specific guides providing valuable insights into structuring these intricate systems effectively and managing the communication and coordination between the different agents within the multi-agent architecture.<sup>8</sup> The ADK documentation encourages developers to thoroughly explore the full set of guides that are available to them.<sup>5</sup> By doing so, developers can gain a deeper and more comprehensive understanding of the framework's capabilities, learn about advanced techniques and best practices, and ultimately be better equipped to build powerful, reliable, and responsible AI agents that meet their specific application requirements.<sup>5</sup> The inclusion of a dedicated guide on "Responsible Agents" within the ADK documentation clearly highlights Google's strong commitment to promoting the ethical and safe development and deployment of AI agents using their framework. By providing specific and actionable guidance on responsible AI practices, the ADK encourages developers to thoughtfully consider the broader societal implications of the AI applications they build and equips them with the necessary knowledge and tools to proactively mitigate potential risks and ensure their agents are used in a beneficial and ethical manner.

## 13. Contributing to the ADK Open Source Project:

### 13.1 Contributing Guide:

As an open-source project, the ADK actively encourages developers from the wider community to contribute to its ongoing development and improvement.<sup>7</sup> To facilitate this collaboration, the official ADK documentation likely includes a comprehensive "Contributing Guide" that clearly outlines the process for getting involved in the project.<sup>7</sup> The GitHub repository for the ADK also hosts a "Code of Conduct," which sets forth the expected

standards for community participation, ensuring a welcoming, inclusive, and respectful environment for all contributors.<sup>5</sup> Furthermore, the documentation for the closely related Agent2Agent (A2A) protocol explicitly mentions clear and well-defined pathways for community contributions<sup>16</sup>, indicating that the ADK project similarly values and encourages involvement from the developer community. The contributing guide would typically provide detailed information and step-by-step instructions on how to report any bugs or issues that are encountered, how to propose new features or enhancements to the framework, and how to submit code contributions following the project's established guidelines.<sup>7</sup> It might also include specific guidelines on coding style and conventions, the required testing procedures for any submitted code, and an overview of the overall development workflow of the ADK project, helping new contributors to integrate smoothly into the development process.<sup>7</sup> The open-source nature of the ADK, coupled with the presence of clear and comprehensive contribution guidelines, signifies a strong commitment to a model of community-driven development and continuous innovation within the rapidly evolving field of AI agents. By opening up the ADK to contributions from the global developer community, Google fosters a collaborative environment where a diverse range of perspectives, expertise, and innovative ideas can contribute to the framework's growth, robustness, and versatility. This collaborative approach can lead to more rapid innovation, a more comprehensive and well-tested toolkit, and a stronger overall ecosystem around the ADK.

## **14. Navigating the ADK API Reference:**

### **14.1 API Documentation:**

A critical and indispensable resource for developers who are working with the ADK is the comprehensive API Reference.<sup>4</sup> This section of the official documentation serves as the definitive source of technical information for the entire ADK framework, providing detailed specifications and explanations for all the classes, methods, functions, and modules that constitute the toolkit.<sup>4</sup> The API reference acts as the primary point of technical reference for developers, offering in-depth details about the purpose of each component within the ADK, the parameters that each function or method accepts, the type of values that are returned, and any potential exceptions or errors that might occur during their use.<sup>4</sup> Developers can explore the structure and organization of the API reference to gain a thorough understanding of the different parts of the ADK framework and how these individual components can be effectively used together to build and extensively customize their AI agents and multi-agent systems.<sup>4</sup> The API reference is designed to allow developers to quickly and efficiently find specific information about any particular ADK component they are interested in, enabling them to effectively utilize the framework's full range of capabilities and troubleshoot any issues or challenges they might encounter during the development process.<sup>4</sup> The ADK documentation strongly encourages developers to consult the API Reference frequently as they work with the framework.<sup>5</sup> This regular consultation will help them to gain a deep and thorough understanding of the ADK's technical details, ensuring they can leverage its features effectively and build robust and sophisticated AI agent applications.<sup>5</sup> A well-structured and comprehensive API reference is absolutely essential for the usability and successful adoption of any software development kit. The ADK's provision of such detailed and readily accessible



documentation empowers developers with the precise technical information they need to understand how to interact with the framework's various components and to build their AI agent applications in an informed and effective manner.

## **15. Exploring the ADK Python GitHub Repository:**

### **15.1 Repository Structure:**

The google/adk-python repository hosted on GitHub serves as the central hub for the ADK's source code and all related resources.<sup>5</sup> Exploring the organization and structure of this repository provides valuable insights into how the codebase is organized and the different modules and packages that collectively form the framework.<sup>5</sup> Key directories within the repository might include the core source files that implement the ADK's functionalities, a collection of illustrative code examples demonstrating various features, a suite of tests designed to ensure the framework's stability and correctness, and the source files for the project's documentation.<sup>5</sup>

### **15.2 Code Examples:**

The GitHub repository contains a variety of practical code examples that clearly demonstrate how to effectively use different features and functionalities of the ADK.<sup>3</sup> These examples can range from very basic demonstrations of how to create a simple AI agent to more complex implementations of multi-agent systems and integrations with various external tools and services.<sup>3</sup> For developers who are new to the ADK, analyzing these provided examples is an excellent way to learn how to apply the framework in real-world scenarios and to understand the recommended best practices for building AI agents using the toolkit.<sup>3</sup> The initial quickstart example that is often provided in the repository's README file offers a first, hands-on introduction to the process of defining and running a basic AI agent using the ADK.<sup>5</sup>

### **15.3 Recent Updates and Discussions:**

Actively checking the commit history of the GitHub repository allows developers to stay well-informed about the most recent changes, bug fixes that have been implemented, and any new features or enhancements that have been added to the ADK.<sup>5</sup> The "Issues" tab on GitHub provides a dedicated platform for developers to report any bugs they encounter while using the ADK, to suggest potential enhancements or new features they believe would be valuable, and to ask questions related to the framework's functionality or usage.<sup>5</sup> Regularly reviewing the issues that have been reported and the discussions surrounding them can provide valuable insights into common problems that other developers are facing and the ongoing conversations within the ADK community.<sup>5</sup> The "Pull Requests" tab on the repository shows proposed changes to the ADK's codebase that have been submitted by contributors from the community.<sup>5</sup> Examining these pull requests can offer a glimpse into the future direction of the project and how the community is actively involved in its development and evolution.<sup>5</sup> It is also important for developers to monitor the repository for official announcements and releases of new versions of the ADK.<sup>5</sup> Staying up-to-date with the latest releases ensures that developers can take advantage of the newest features, performance improvements, and bug fixes.<sup>5</sup> The level of activity on the repository, such as the number of stars it has received, the number of developers watching it, and the number of times it has

been forked, can also provide an indication of the broader community's interest in and engagement with the ADK project.<sup>5</sup> The video mentioned by the user also specifically directs developers to the GitHub repository as the primary starting point for anyone interested in getting started with the ADK and exploring its capabilities.<sup>18</sup> The repository's content reinforces the ADK's fundamental "code-first" development approach, where the behavior and underlying logic of AI agents are defined directly through Python code.<sup>5</sup> The repository likely contains practical examples of simple yet functional multi-agent systems that have been built using Python and the ADK, further illustrating the framework's capabilities and ease of use in creating collaborative AI applications.<sup>3</sup> The ADK's GitHub repository serves as an absolutely vital resource for developers who are working with the framework. It provides direct access to the source code, a wealth of illustrative examples that demonstrate how to use the ADK in practice, and a central platform for community interaction, allowing developers to stay informed about the project's ongoing development, contribute their own improvements, and learn from the experiences of others.

16. Conclusion:

The Google Agent Development Kit (ADK) stands out as a powerful and highly versatile framework for building the next generation of AI-powered applications, with a clear and significant emphasis on facilitating the development of sophisticated multi-agent systems. Its core strengths lie in the remarkable flexibility and fine-grained control it offers to developers through its developer-centric "code-first" approach, its seamless and deep integration with the robust Google Cloud ecosystem and a wide array of cutting-edge language models, and its rich and extensive set of tools that allow for significant extension of agent capabilities. The ADK's inherent support for crucial aspects of modern AI development, such as real-time streaming for more natural interactions, comprehensive evaluation tools for ensuring performance and reliability, and dedicated guidance for responsible AI development, further positions it as a comprehensive and forward-thinking solution for building robust and ethical AI agent applications. As an open-source project, the ADK benefits immensely from the active contributions of a vibrant developer community, fostering a collaborative environment that encourages continuous innovation and growth within the rapidly evolving field of artificial intelligence agents. With its comprehensive feature set, its strong backing from Google's advanced AI research and infrastructure, and its commitment to open collaboration, the ADK holds significant potential to empower developers in creating truly innovative and impactful AI applications that can leverage the collective intelligence and coordinated efforts of multiple interacting agents to solve complex real-world problems.

Key Tables:

- **Table 1: Compatible Language Models and Integration Methods (Section 4)**

Language Model	Provider	Integration Method	Key Features/Considerations
Gemini	Google	Native, Vertex AI	Optimized

			performance, advanced reasoning, multimodal capabilities
Models in Model Garden	Google & Others	Vertex AI Model Garden	Wide selection of pre-trained models, managed platform
Various LLMs	Anthropic, Meta, Mistral, AI21	LiteLLM	Unified API for multiple providers, easy switching between models
Popular LLMs	Various	Direct API Integration, LiteLLM	Flexibility in choosing models based on specific needs

• **Table 2: Types of Tools in ADK (Section 5)**

Tool Type	Description	Examples/Integration Methods
Function Tools	Custom Python functions to extend agent capabilities.	Interacting with APIs, performing data processing, executing custom logic.
Built-in Tools	Pre-packaged tools for common tasks.	Web searching (google_search), code execution, retrieving weather/time.
Third-Party Tools	Integration with external libraries and services.	LangChain, CrewAI, LlamaIndex.
Google Cloud Tools	Leveraging specific services within the Google Cloud platform.	Vertex AI for model deployment, connecting to Google Cloud APIs.
MCP Tools	Connecting to diverse data sources and capabilities using Model Context Protocol.	Accessing data through MCP-compatible tools, efficient reasoning over big data.

OpenAPI Tools	Integrating with services that expose an OpenAPI specification.	Interacting with a wide range of web services in a standardized way.
---------------	---	--

## Works cited

1. Agent Development Kit: Making it easy to build multi-agent applications, accessed on April 10, 2025, <https://developers.googleblog.com/en/agent-development-kit-easy-to-build-multi-agent-applications/>
2. Google's Agent Development Kit (ADK): Revolutionizing Multi-Agent Application Development | by Zen Solutions | Apr, 2025 | Medium, accessed on April 10, 2025, [https://medium.com/@hams\\_ollo/googles-agent-development-kit-adk-revolutionizing-multi-agent-application-development-f74f9a08cdce](https://medium.com/@hams_ollo/googles-agent-development-kit-adk-revolutionizing-multi-agent-application-development-f74f9a08cdce)
3. Google Releases Agent Development Kit (ADK): An Open-Source AI Framework Integrated with Gemini to Build, Manage, Evaluate and Deploy Multi Agents - MarkTechPost, accessed on April 10, 2025, <https://www.marktechpost.com/2025/04/09/google-releases-agent-development-kit-adk-an-open-source-ai-framework-integrated-with-gemini-to-build-manage-evaluate-and-deploy-multi-agents/>
4. Agent Development Kit - Google, accessed on April 10, 2025, <https://google.github.io/adk-docs/>
5. google/adk-python: An open-source, code-first Python toolkit for building, evaluating, and deploying sophisticated AI agents with flexibility and control. - GitHub, accessed on April 10, 2025, <https://github.com/google/adk-python>
6. Build and manage multi-system agents with Vertex AI | Google Cloud Blog, accessed on April 10, 2025, <https://cloud.google.com/blog/products/ai-machine-learning/build-and-manage-multi-system-agents-with-vertex-ai>
7. Get Started - Agent Development Kit - Google, accessed on April 10, 2025, <https://google.github.io/adk-docs/get-started/>
8. Tutorial - Agent Development Kit - Google, accessed on April 10, 2025, <https://google.github.io/adk-docs/get-started/tutorial/>
9. Google Cloud Preps for Agentic AI Era with 'Ironwood' TPU, New Models and Software, accessed on April 10, 2025, <https://www.hpcwire.com/2025/04/09/google-cloud-preps-for-agentic-ai-era-with-ironwood-tpu-new-models-and-software/>
10. Quickstart: Build an agent with the Agent Development Kit | Generative AI on Vertex AI, accessed on April 10, 2025, <https://cloud.google.com/vertex-ai/generative-ai/docs/agent-development-kit/quickstart>
11. Quickstart - Agent Development Kit - Google, accessed on April 10, 2025, <https://google.github.io/adk-docs/get-started/quickstart/>
12. accessed on January 1, 1970,

- [https://google.github.io/adk-docs/get-started/quickstart\\_streaming/](https://google.github.io/adk-docs/get-started/quickstart_streaming/)
13. Google Doubles Down on Developer AI with New Tools, Agent Ecosystem, and Hyperscale Infrastructure -- ADTmag, accessed on April 10, 2025, <https://adtmag.com/articles/2025/04/09/google-doubles-down-on-developer-ai-with-new-tools.aspx>
  14. Google goes all in on agent development, Gemini at Google Cloud Next 25 - SD Times, accessed on April 10, 2025, <https://sdtimes.com/ai/google-goes-all-in-on-agent-development-gemini-advancements-at-google-cloud-next-25/>
  15. Google Cloud sees multi-agent AI systems as 'next frontier' - Fierce Healthcare, accessed on April 10, 2025, <https://www.fiercehealthcare.com/ai-and-machine-learning/google-cloud-builds-out-ai-agent-capabilities-healthcare-highmark-health>
  16. Announcing the Agent2Agent Protocol (A2A) - Google for Developers Blog, accessed on April 10, 2025, <https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/>
  17. Responsible Agents - Agent Development Kit - Google, accessed on April 10, 2025, <https://google.github.io/adk-docs/guides/responsible-agents/>
  18. Google Launches an Agent SDK - Agent Development Kit - YouTube, accessed on April 10, 2025, <https://www.youtube.com/watch?v=G9wnpfW6lZY>